




A local, agentic tool that reads a CV and a job offer and tells you — explainably — how well they fit, then rewrites the CV to fit better. It works from both sides of the table: the **Applier** (job-seeker) and the **Recruiter** (hiring).

Everything runs on your own machine and on **your own AI account**. Nothing is stored; nothing leaves your computer except the calls you make to your AI and to the job boards you switch on. It is not a website you log into — it's an app you run.

01 THE MACHINE AT A GLANCE



One AI door (two swappable backends) · a deterministic core that needs no AI · an office of agents, two of which run in a quality loop.

02 THE TWO MODES

Applier **JOB-SEEKER**

- **Analyze** — an explainable compatibility score: ATS keyword coverage, recruiter gating (languages, years), and the AI's semantic read, each shown separately so you see *why*.
- **Find offers** across free job boards, or paste one link.
- **Regenerate** — rewrites your CV for a specific offer with a 1–100 aggressiveness slider, then hands you an ATS-safe PDF.

Recruiter **HIRING**

- **Screen & rank** a batch of candidate CVs against one role.
- **AI-writing signals** — heuristics + a reading model flag likely AI-written CVs (as *signals*, never *proof*).
- **Verify sources** — checks that links and GitHub repos are real.
- **Source candidates** legally — GitHub search + a Google/LinkedIn X-ray string you run yourself.

03 THE OFFICE OF AGENTS

AGENT	SIDE	WHAT IT DOES
Analyst	both	Reads the offer + CV, extracts the requirements, and produces the explainable report.
Tailor	applier	Rewrites the CV at the slider's aggressiveness. Anything it adds beyond your real CV is written to the truth-ledger.
Critic	applier	Scores each rewrite and enforces the gate: any un-ledgered claim fails it.
Examiner	recruiter	Judges AI-writing likelihood (Low / Medium / High) with evidence — a signal, not a verdict.
Verifier	recruiter	Confirms links and GitHub repositories exist (owner, stars, last commit).
Sourcer	recruiter	Turns a role into a GitHub search and an X-ray string for consent-based sourcing.

04 THE QUALITY LOOP THE HEART OF IT

The Applier's rewrite is not one shot. The **Tailor** writes a version; the **Critic** grades it against a rubric — keyword coverage, readability, ATS-parseability, and **truthfulness** as a hard gate. If it doesn't pass (target ≥ 85), the Critic's feedback goes back to the Tailor and it tries again — up to three passes, keeping the best. The same philosophy as the rest of Gradian's factory: nothing ships below the bar.

Truth-locked. The slider can be aggressive, but every claim it adds beyond your source CV lands in the **truth-ledger** — "⚠ verify before sending" — so you never walk into an interview not knowing what your own CV says.

05 HOW YOU CONNECT YOUR AI — TWO WAYS

Claude Code DEFAULT · NO KEY

Install Claude Code and sign in. Gradian Match calls it under the hood on your subscription. Nothing to buy, nothing to paste.

Anthropic API key FASTER

Put `ANTHROPIC_API_KEY` in `.env`. This talks to the API directly — lighter and faster per call than the CLI, ideal for the multi-call rewrite loop.

Either way it is **local and self-hosted**. There is deliberately **no hosted version**: CVs are personal data, and keeping them on your machine is the whole point. The tool uses **no MCP servers** — it reaches the AI directly through the CLI or the SDK.

06 PRIVACY & HONESTY

Everything is in-memory per session; temporary files are deleted; the server binds to `127.0.0.1` only and rejects cross-origin requests. In recruiter mode you process third parties' data — the legal basis is yours, and the tool holds nothing. The AI-signal flags are explicitly **signals, not proof**: don't auto-reject anyone on them.